

Textbook: William Stallings, Data and Computer Communications

Data Communications and Networking

Chapter 17

Transport Protocols

<http://cs.nju.edu.cn/yangxc>
xcy@nju.edu.cn




Transport Protocol


- Sits above a network or internetwork, which provides network-related services
- Provides transport services (TS) to upper layer users such as FTP, SMTP, and TELNET
- Uses some lower layer services such as IP protocol to support communication between local and remote transport entities



An Internetwork

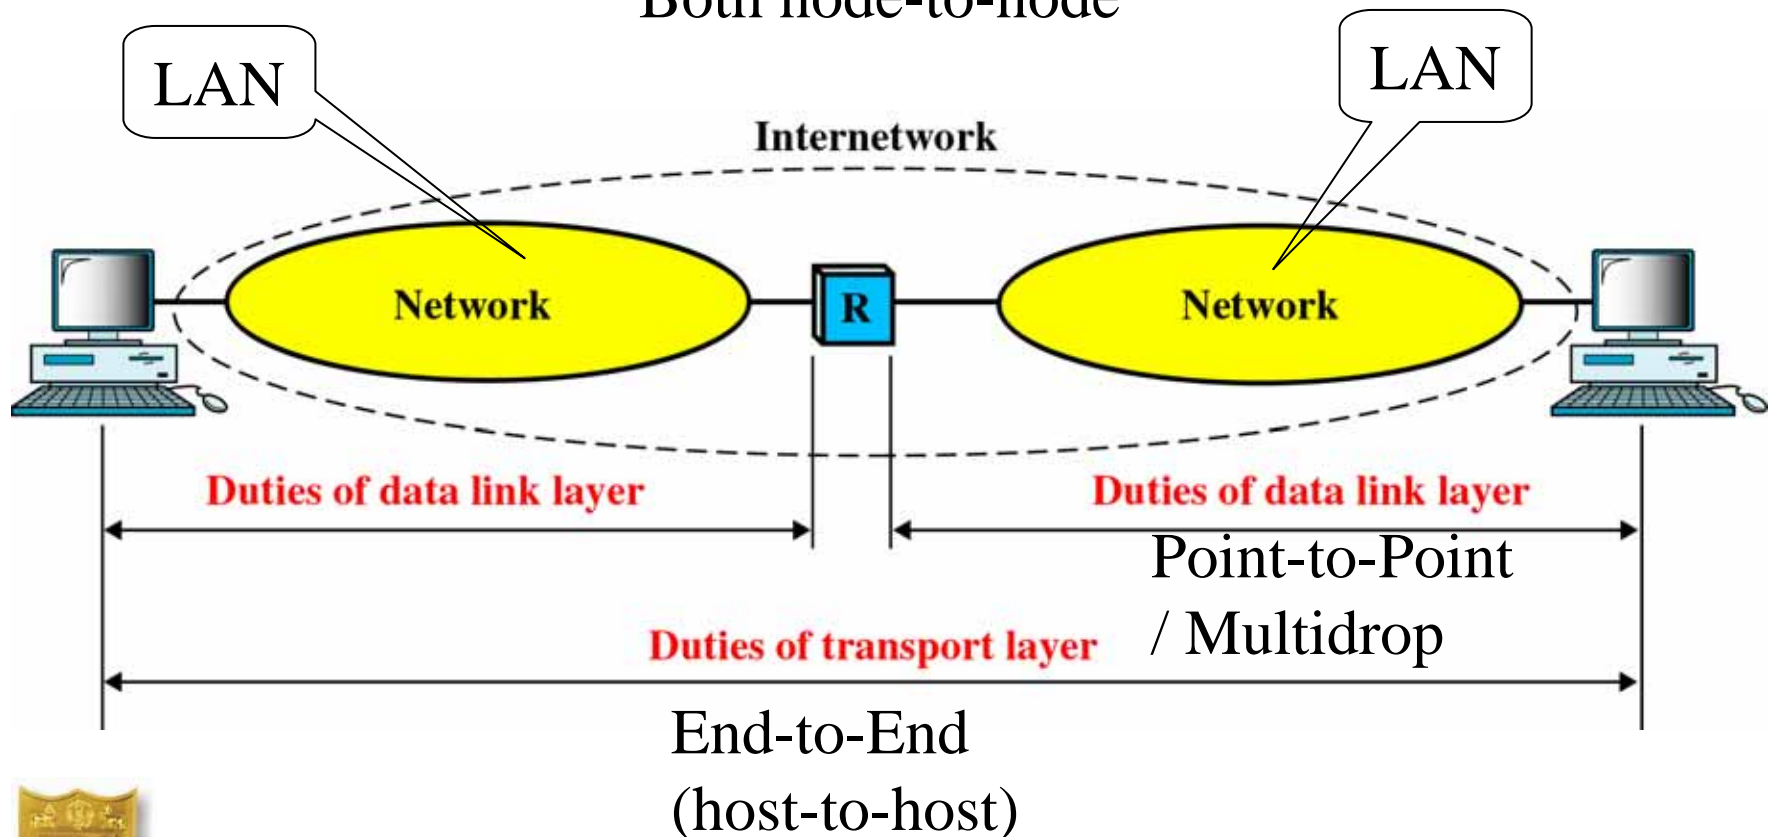


Transport Layer Concept



Transport Layer Compared with Data Link Layer

Both node-to-node

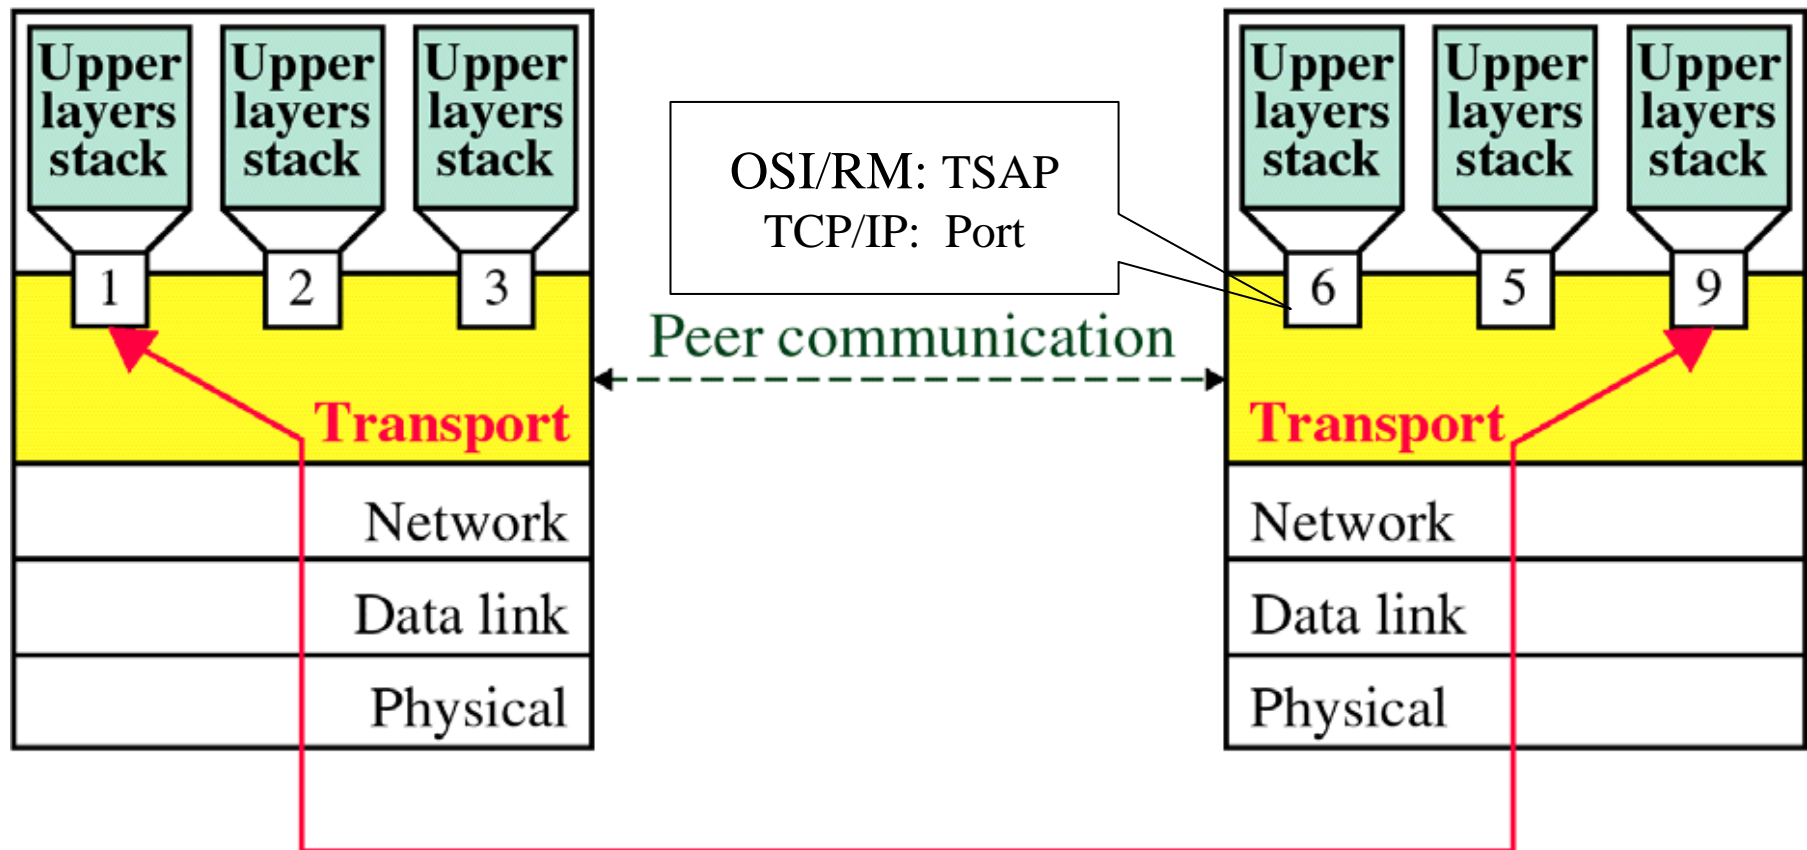


Transport Layer Compared with Network or Internetwork Layer

- End-to-end Delivery
 - Network or internetwork layer
 - Oversees end-to-end delivery of individual packets
 - No see any relationship between packets
 - Deliver packets to host
 - Transport layer
 - Make sure entire message (not just a single packet) arrive intact
 - Oversee end-to-end delivery of an entire message
 - Deliver messages to application processes in host




Service Points



TSAP - Transport Service Access Point



Transport Layer Duties

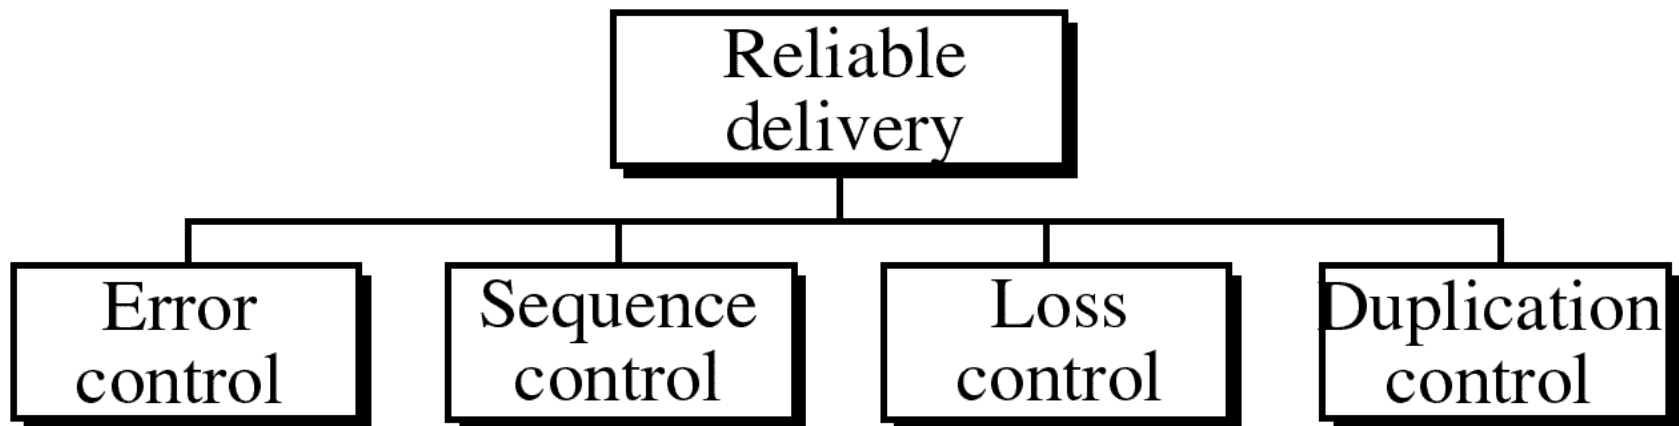


Addressing

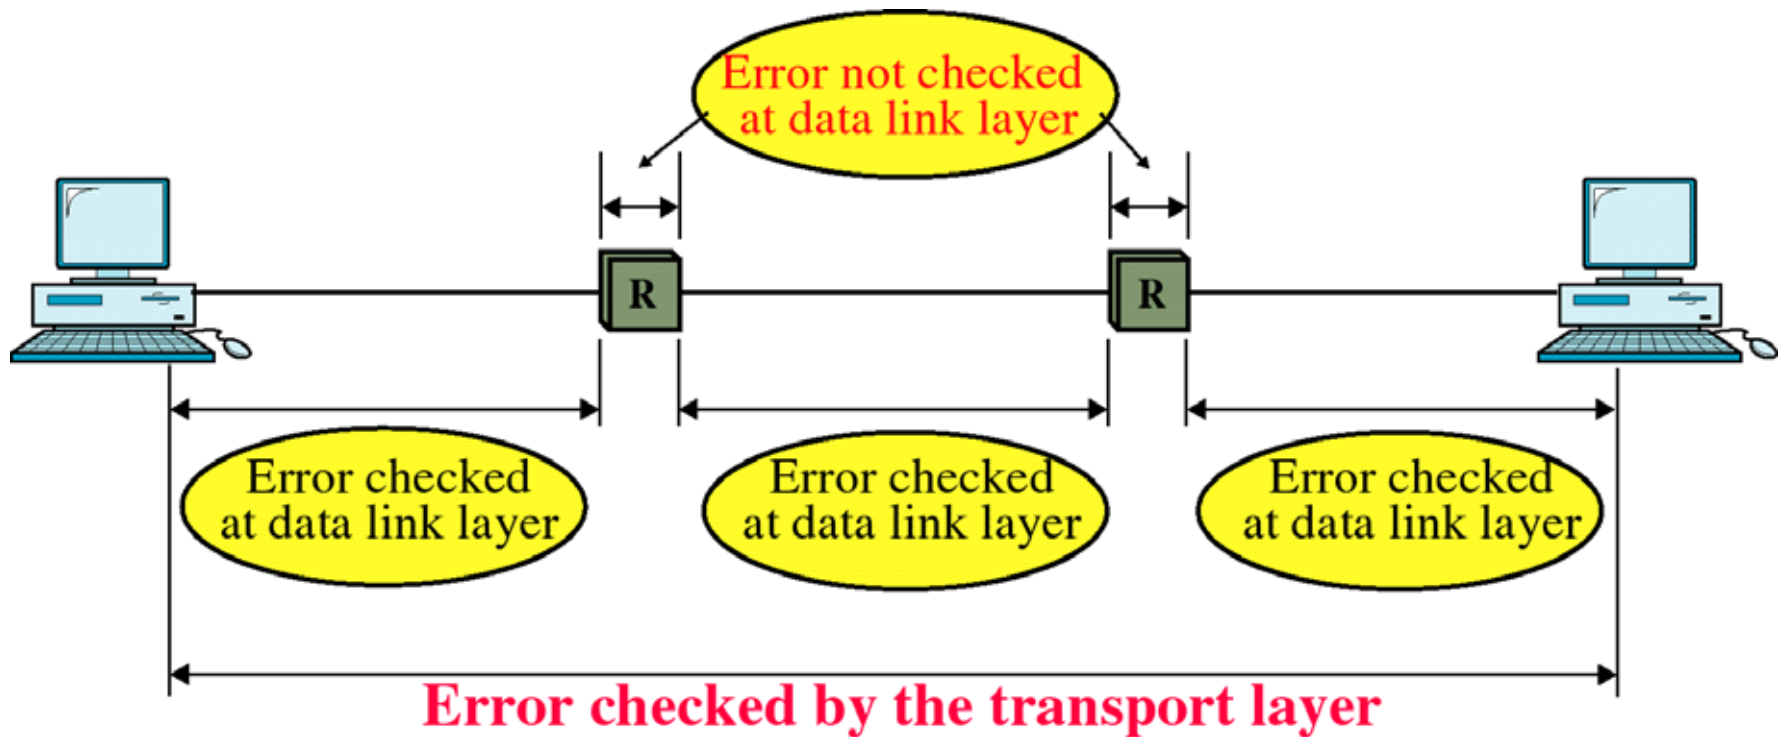
- To ensure accurate delivery from service point to service point, another level of addressing will be needed in addition to those at data link and network levels
 - **Data link level protocols** need to know which two computers within a network are communicating
 - **Network level protocols** need to know which two computers within an internetwork are communicating
 - **Transport level protocol** needs to know which upper layer protocols are communicating



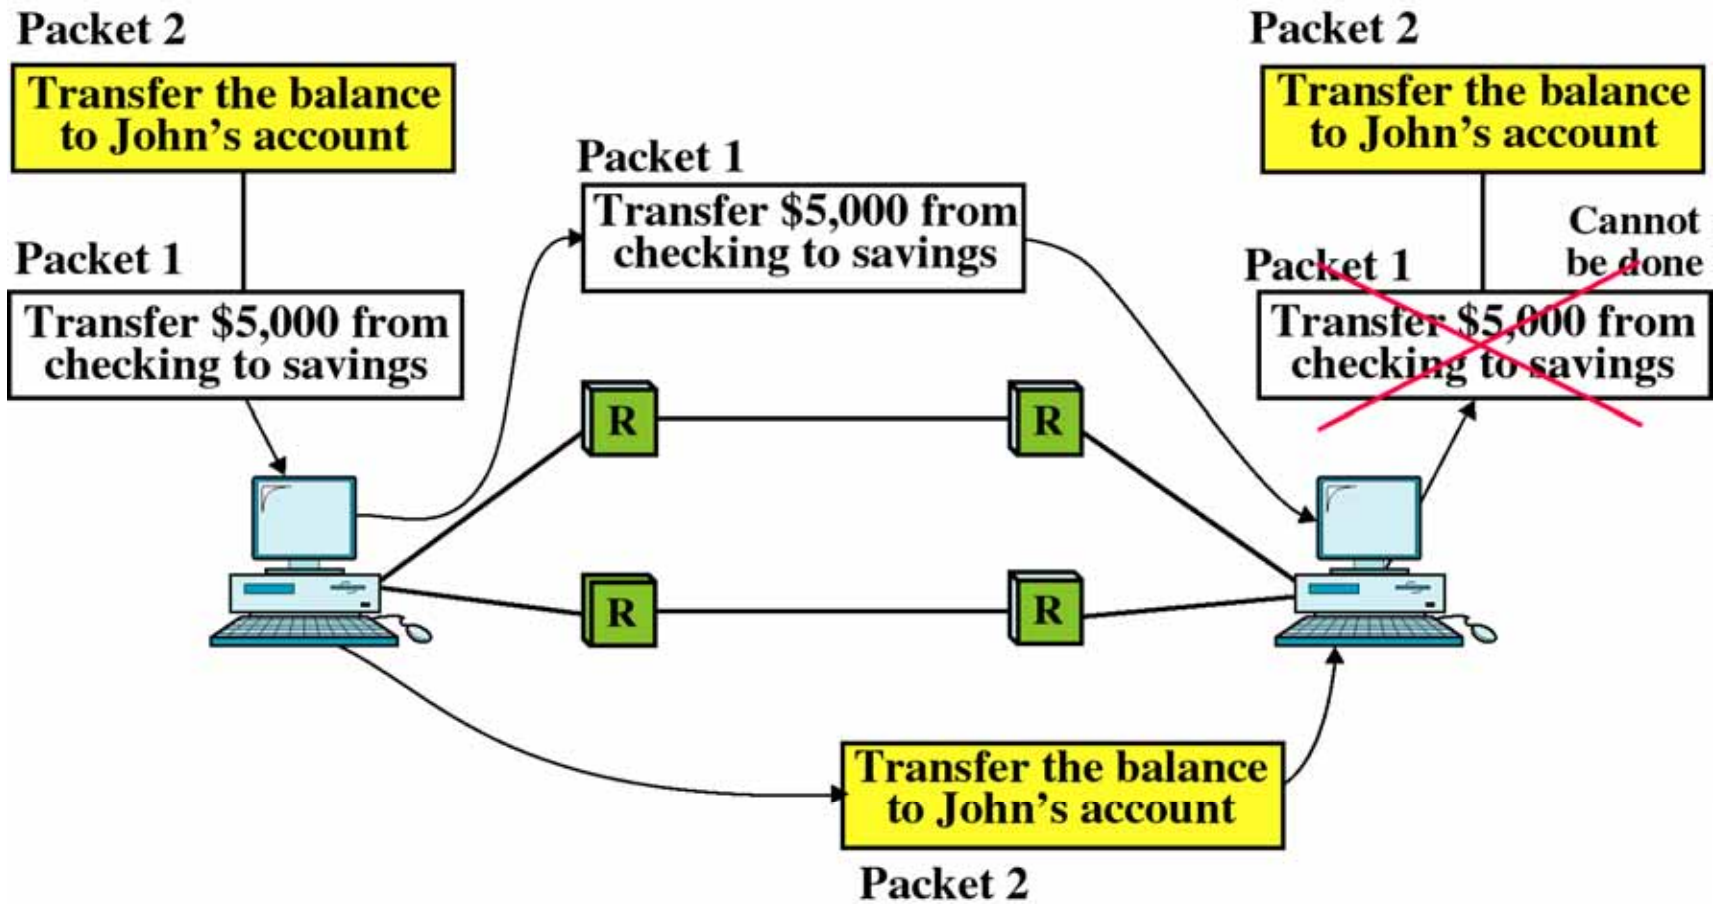
Reliable Delivery




Transport and Data Link Error Control



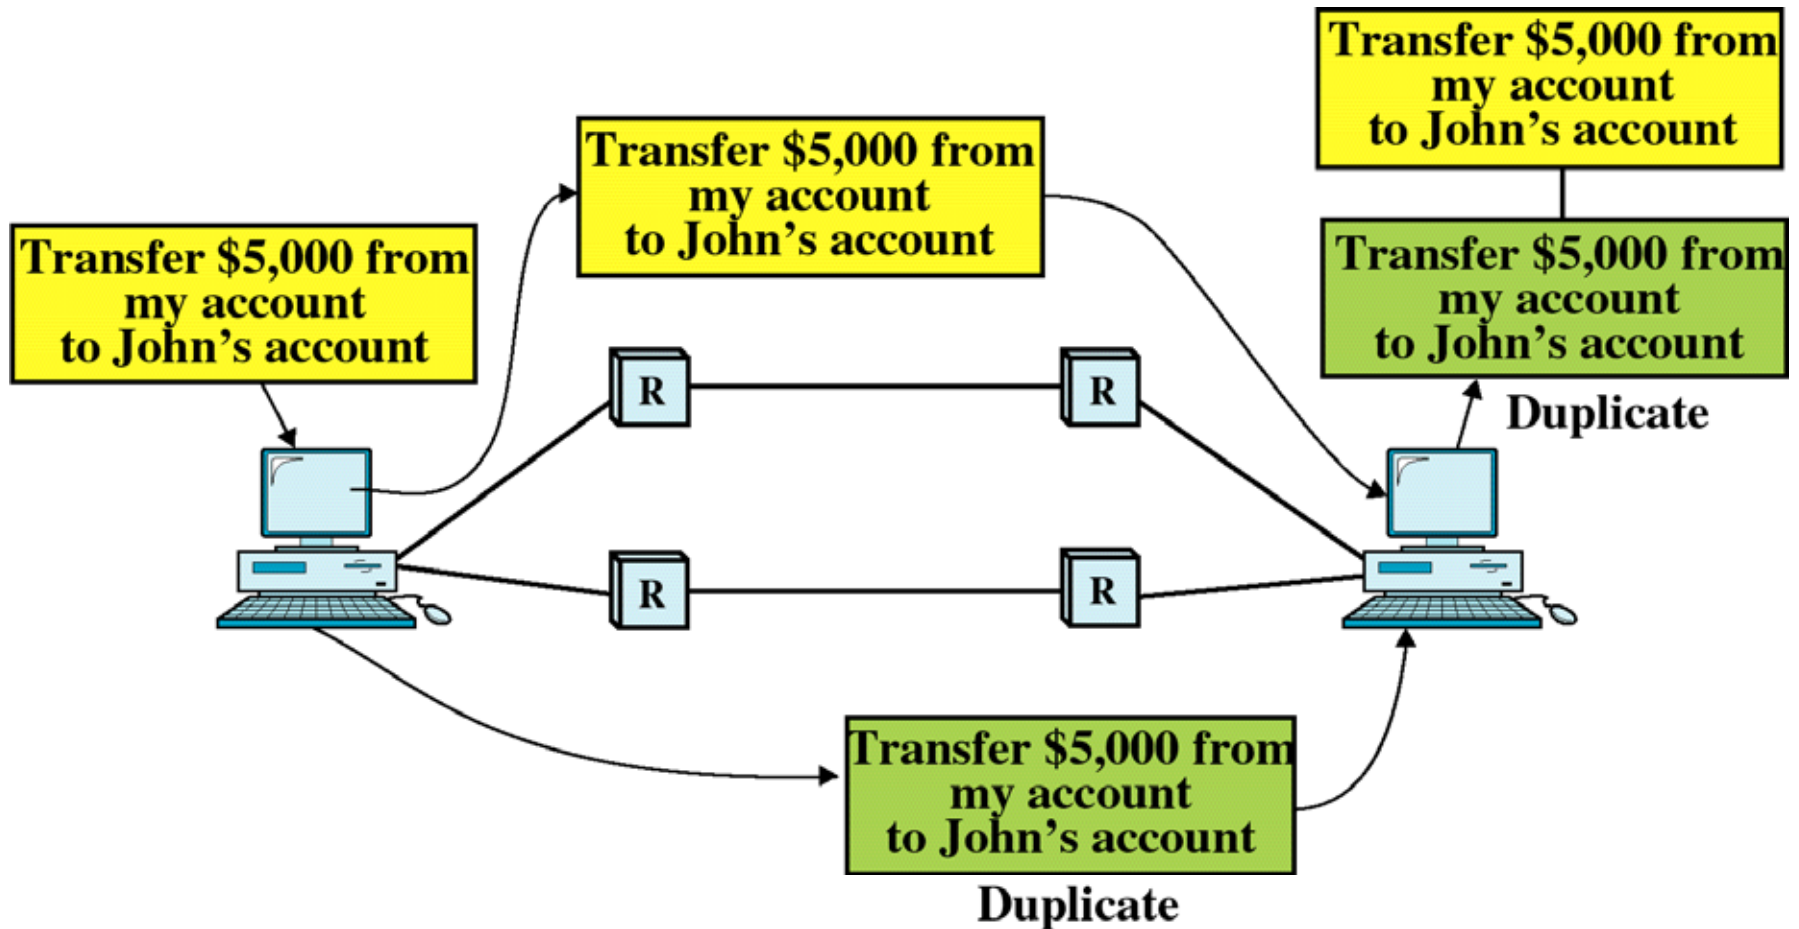
Sequence Control





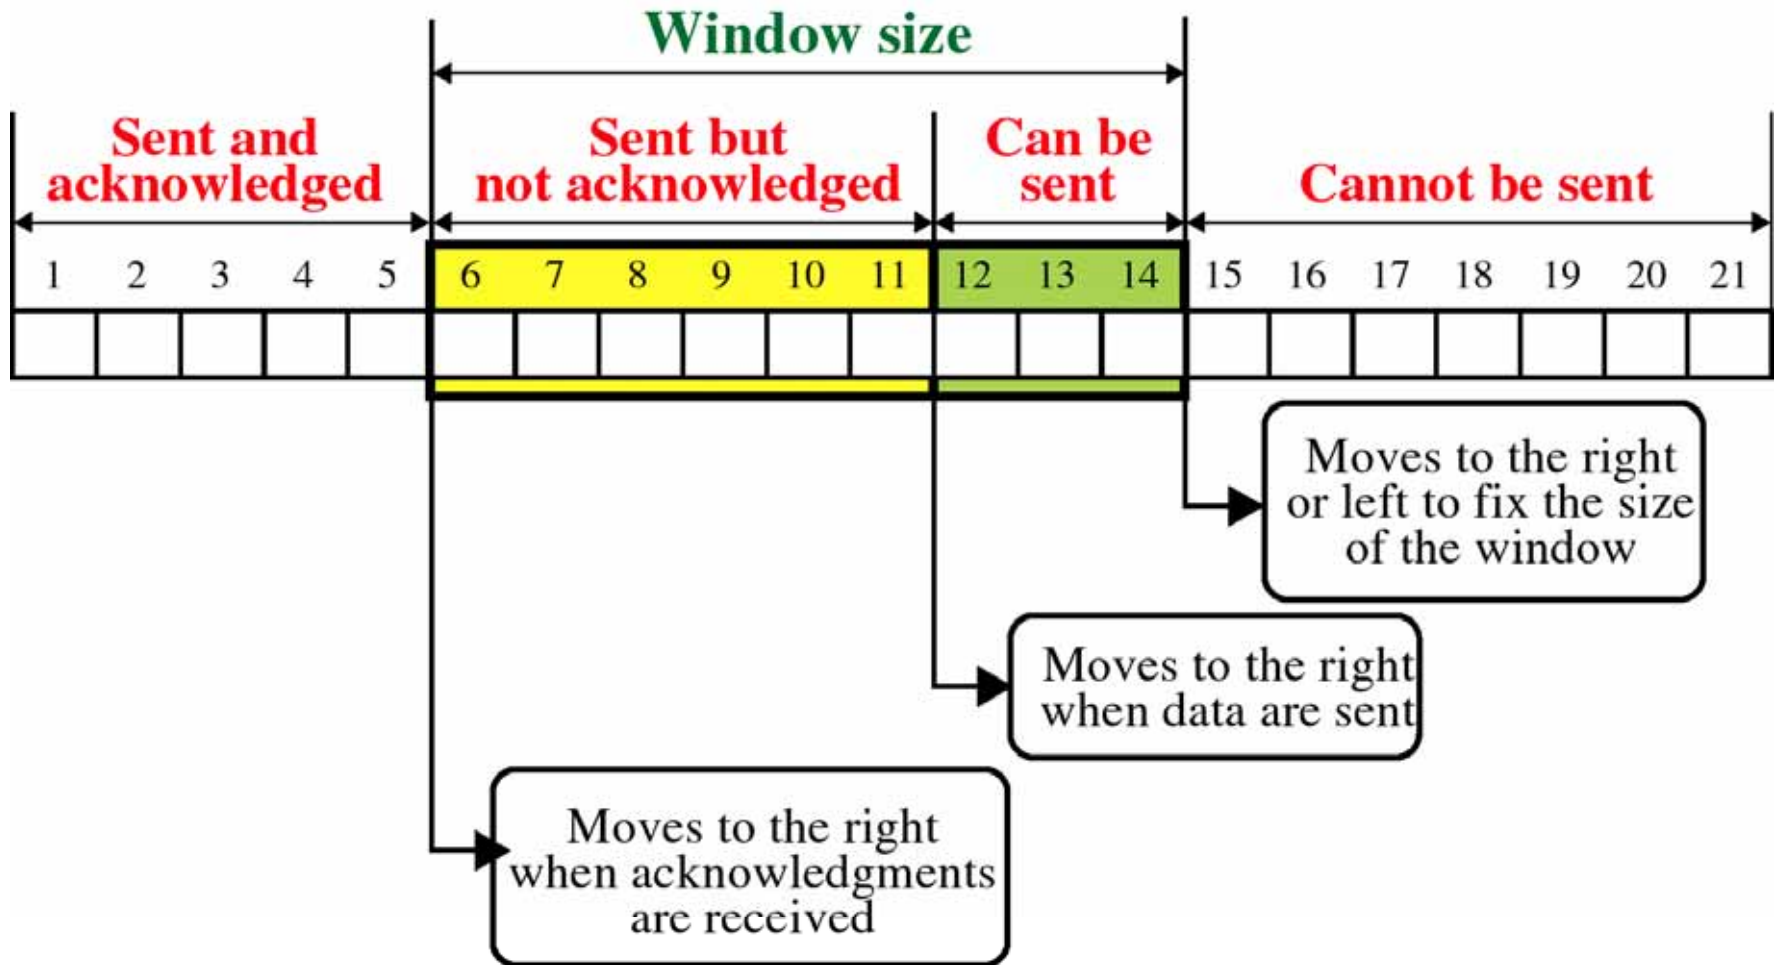
Loss Control





Duplication Control







Flow Control



Multiplexing



Services of Transport Layer and Network Layer



Connection Oriented Transport Protocol Mechanisms

- Logical connection
 - Establishment
 - Maintenance
 - Termination
- Reliable
- e.g. TCP



Reliable Sequencing Network Service

- Assume arbitrary length message
- Assume virtually 100% reliable delivery by network service
- Examples
 - reliable packet switched network using X.25
 - frame relay using LAPF control protocol
 - IEEE 802.3 using connection oriented LLC service
- Transport service is end to end protocol between two systems on same network



Issues in a Simple Transport Protocol

- Addressing
- Multiplexing
- Flow Control
- Connection establishment and termination



Addressing

- Target user specified by:

- **User identification**

- Usually (host, port)
 - Called a socket in TCP
 - Port represents a particular transport service (TS) user

48-bit socket =

32-bit IP address + 16-bit port number

- **Transport entity identification**

- Generally only one per host
 - If more than one, then usually one of each type
 - Specify transport protocol (TCP, UDP)

- **Host address**

- An attached network device
 - In an internet, a global internet address

- **Network number**



Finding Addresses

- Four methods (2 static and 2 dynamic)
 - Know address ahead of time
 - e.g. collection of network device statistics
 - Well known addresses
 - e.g. Time sharing and word procssing
 - Name server
 - Sending process request to well known address



Multiplexing

- Multiple users employ same transport protocol
- User identified by port number or service access point (SAP)
- Upward and downward multiplexing
- May also multiplex with respect to network services used
 - e.g. multiplexing a single virtual X.25 circuit to a number of transport service user
 - X.25 charges per virtual circuit connection time



Flow Control

- Longer transmission delay between transport entities compared with actual transmission time
 - Delay in communication of flow control info
- Variable transmission delay
 - Difficult to use timeouts
- Flow may be controlled because:
 - The receiving user can not keep up
 - The receiving transport entity can not keep up
- Results in buffer filling up



Coping with Flow Control Requirements (1)

- Do nothing
 - Segments that overflow are discarded
 - Sending transport entity will fail to get ACK and will retransmit
 - Thus further adding to incoming data
- Refuse further segments
 - Backpressure
 - Clumsy
 - Multiplexed connections are controlled on aggregate flow



Coping with Flow Control Requirements (2)

- Use fixed sliding window protocol
 - See chapter 7 for operational details
 - Works well on reliable network
 - Failure to receive ACK is taken as flow control indication
 - Does not work well on unreliable network
 - Can not distinguish between lost segment and flow control
- Use credit scheme



Credit Scheme

- Greater control on reliable network
- More effective on unreliable network
- Decouples flow control from ACK
 - May ACK without granting credit and vice versa
- Each octet has sequence number
- Each transport segment has seq number, ack number and window size in header



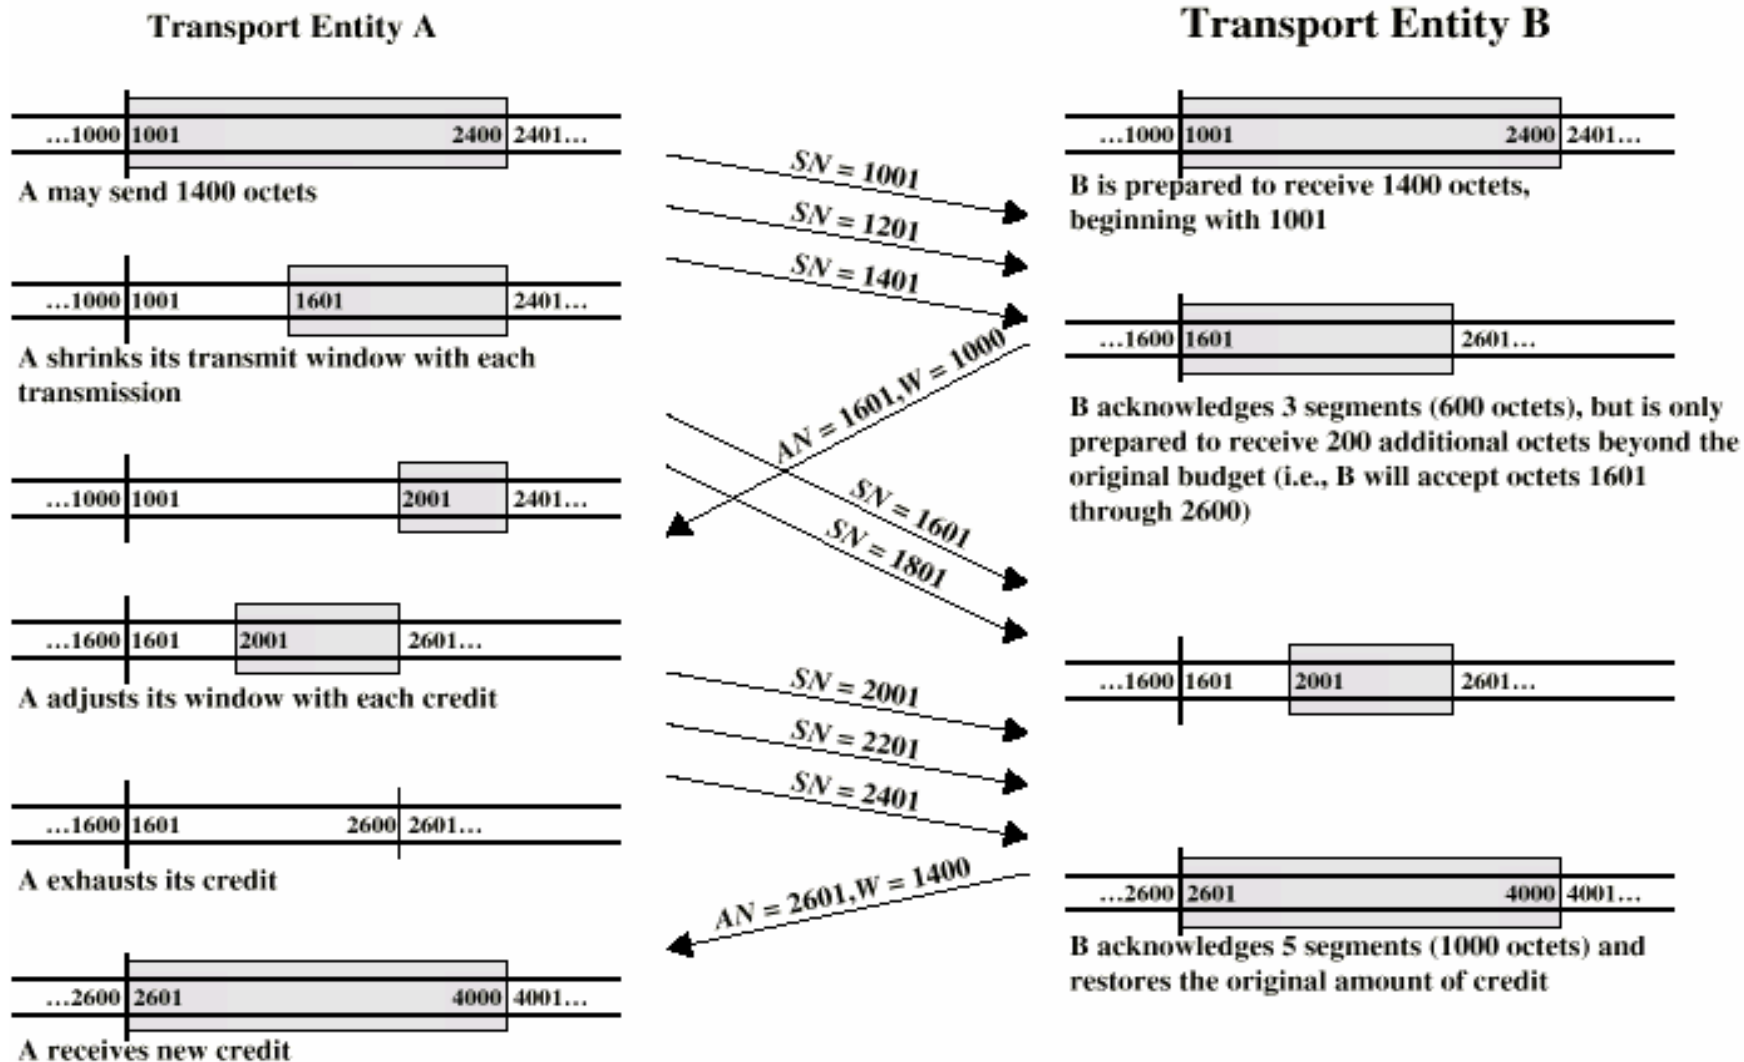
Use of Header Fields

- When sending, seq number is that of first octet in segment
- ACK includes $AN=i$, $W=j$
- All octets through $SN=i-1$ acknowledged
 - Next expected octet is i
- Permission to send additional window of $W=j$ octets
 - i.e. octets through $i+j-1$



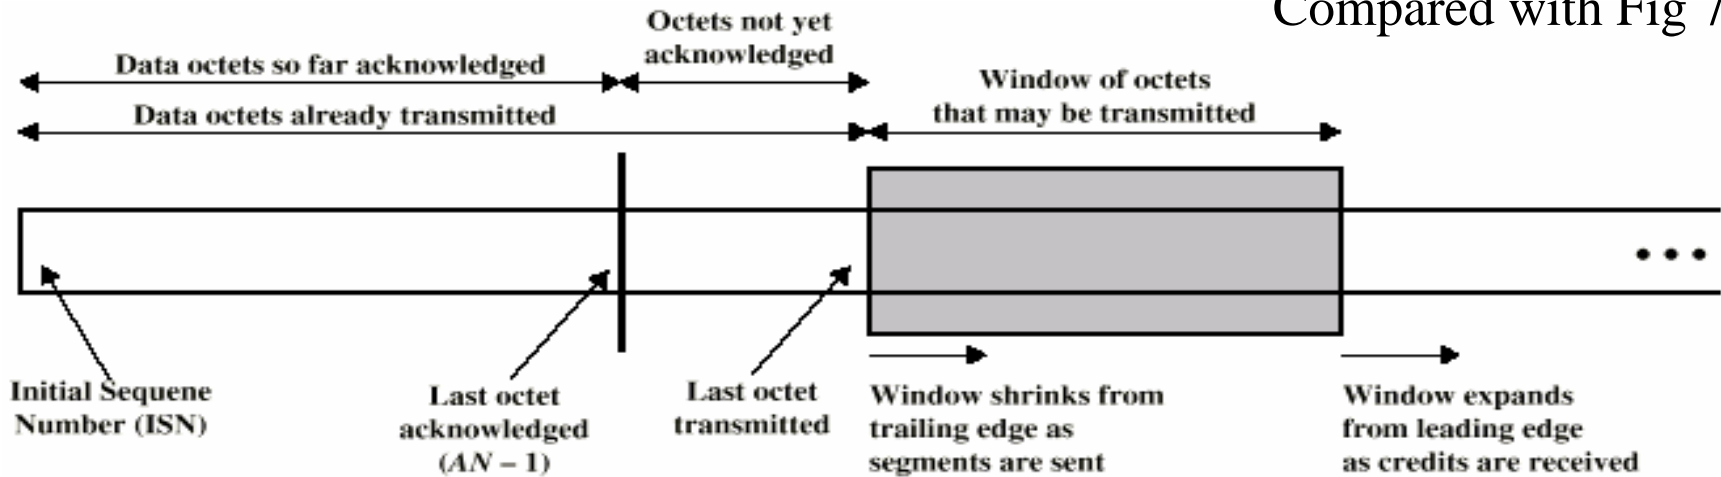
Compared with Fig 7.4

Credit Allocation

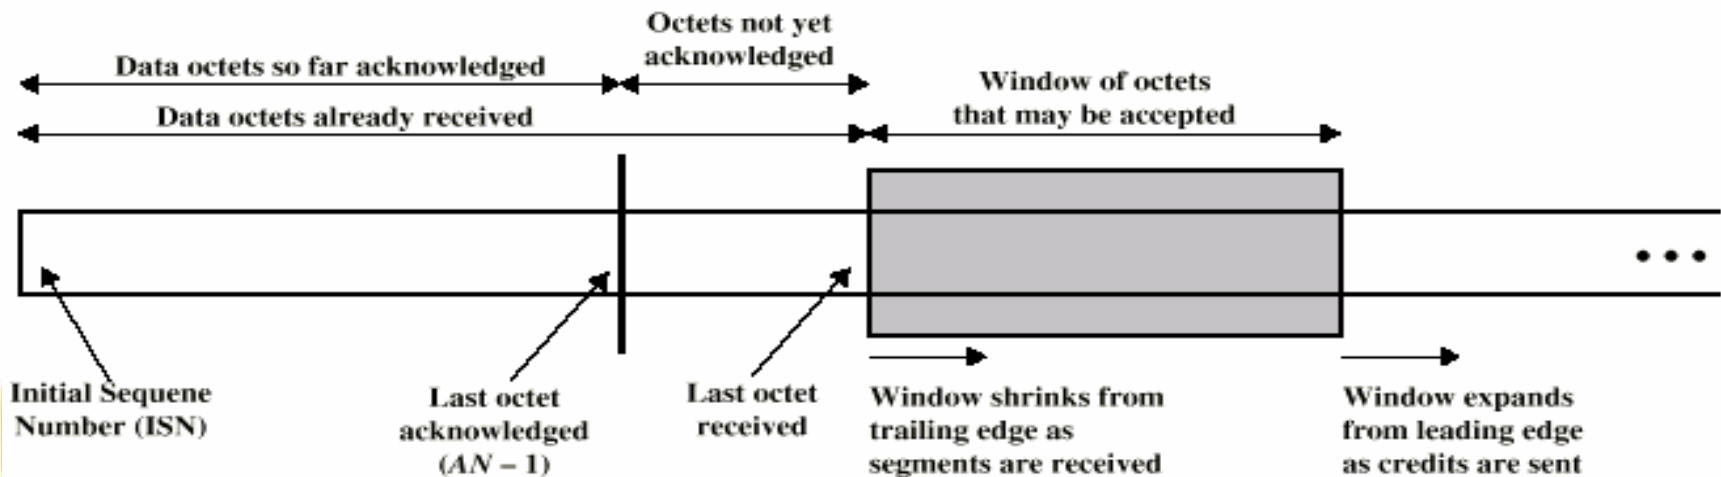


Sending and Receiving Perspectives

Compared with Fig 7.3



(a) Send sequence space




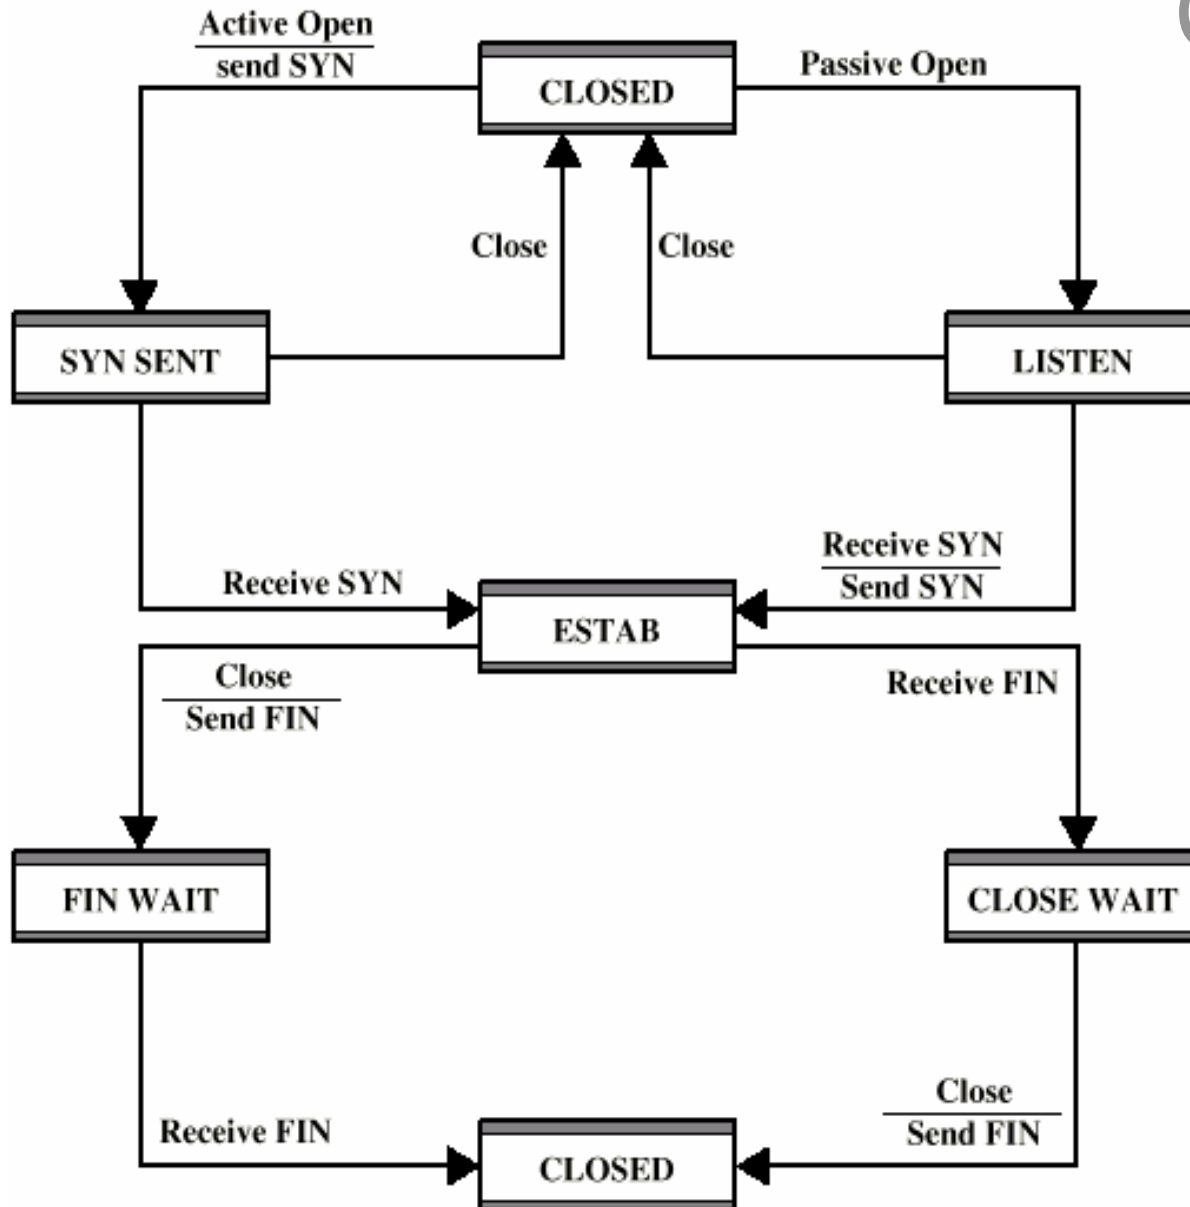
(b) Receive sequence space

Establishment and Termination

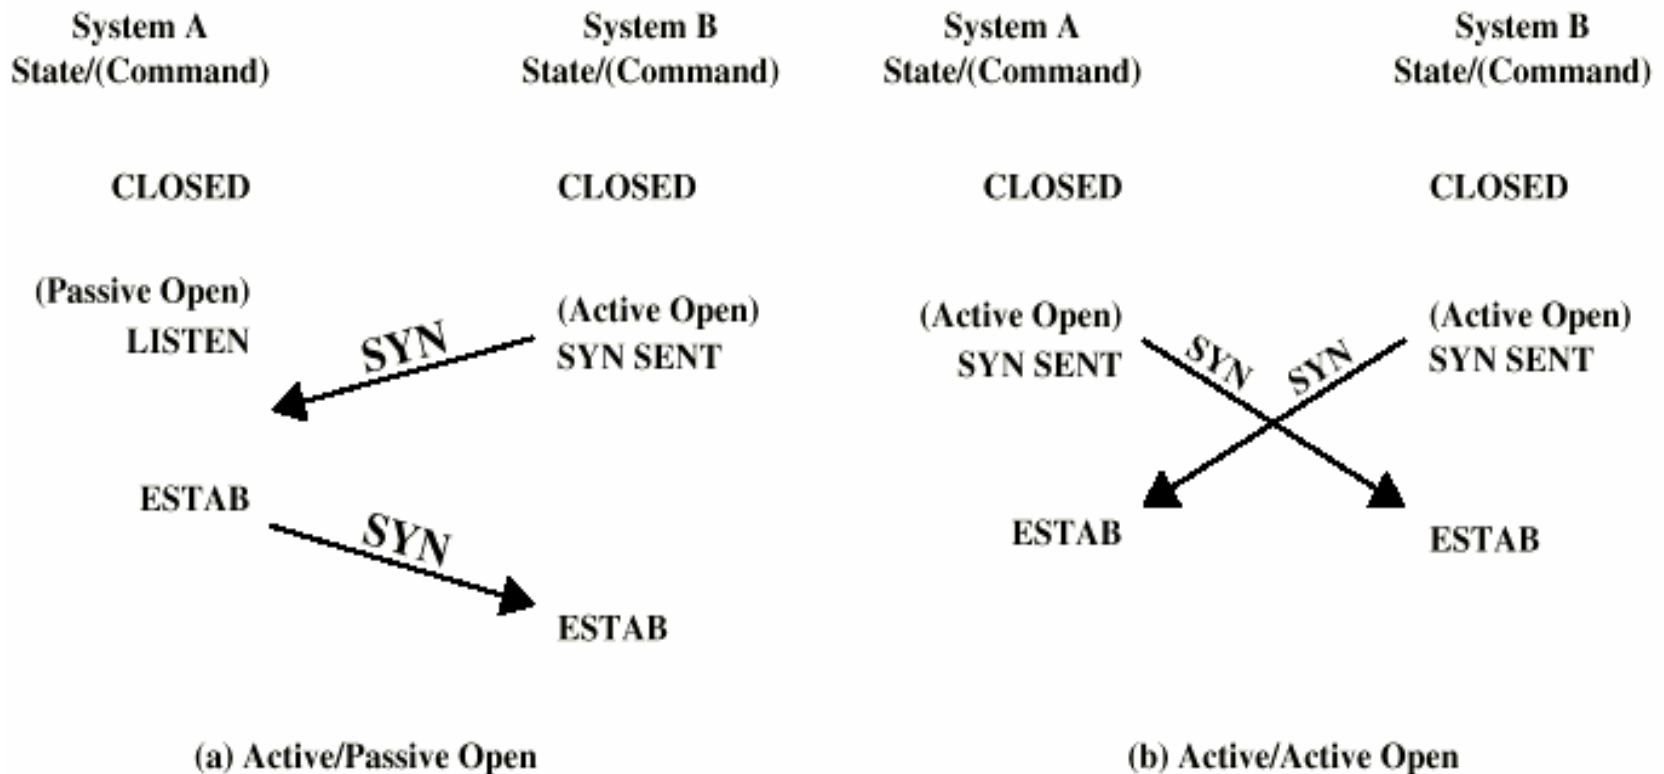
- Allow each end to know the other exists
- Negotiation of optional parameters
 - Max. segment size, Max. window size, QoS
- Triggers allocation of transport entity resources
 - Buffer space, entry in connection table
- By mutual agreement



Connection State Diagram



Connection Establishment



Not Listening

- Reject with RST (Reset)
- Queue request until matching open issued
- Signal TS user to notify of pending request
 - May replace passive open with accept



Termination

- Either or both sides
- By mutual agreement
- Abrupt termination
- Or graceful termination
 - Close wait state must accept incoming data until FIN received



Side Initiating Termination

- TS user Close request
- Transport entity sends FIN, requesting termination
- Connection placed in FIN WAIT state
 - Continue to accept data and deliver data to user
 - Not send any more data
- When FIN received, inform user and close connection



Side Not Initiating Termination

- FIN received
- Inform TS user Place connection in CLOSE WAIT state
 - Continue to accept data from TS user and transmit it
- TS user issues CLOSE primitive
- Transport entity sends FIN
- Connection closed
- All outstanding data is transmitted from both sides
- Both sides agree to terminate



Unreliable Network Service

- Examples
 - internet using IP
 - frame relay using LAPF
 - IEEE 802.3 using unacknowledged connectionless LLC
- Segments may get lost
- Segments may arrive out of order



Problems

- Ordered Delivery
- Retransmission strategy
- Duplication detection
- Flow control
- Connection establishment
- Connection termination
- Crash recovery



Ordered Delivery

- Segments may arrive out of order
- Number segments sequentially
- TCP numbers each octet sequentially
- Segments are numbered by the first octet number in the segment



Retransmission Strategy

- Segment damaged in transit
- Segment fails to arrive
- Transmitter does not know of failure
- Receiver must acknowledge successful receipt
- Use cumulative acknowledgement
- Time out waiting for ACK triggers re-transmission



Timer Value


- Fixed timer
 - Based on understanding of network behavior
 - Can not adapt to changing network conditions
 - Too small leads to unnecessary re-transmissions
 - Too large and response to lost segments is slow
 - Should be a bit longer than round trip time
- Adaptive scheme
 - May not ACK immediately
 - Can not distinguish between ACK of original segment and re-transmitted segment
 - Conditions may change suddenly



Duplication Detection

- If ACK lost, segment is re-transmitted
- Receiver must recognize duplicates
- Duplicate received prior to closing connection
 - Receiver assumes ACK lost and ACKs duplicate
 - Sender must not get confused with multiple ACKs
 - Sequence number space large enough to not cycle within maximum life of segment
- Duplicate received after closing connection





Flow Control


- Credit allocation
- Problem if $AN=i$, $W=0$ closing window
- Send $AN=i$, $W=j$ to reopen, but this is lost
- Sender thinks window is closed, receiver thinks it is open
- Use window timer
- If timer expires, send something
 - Could be re-transmission of previous segment



Connection Establishment

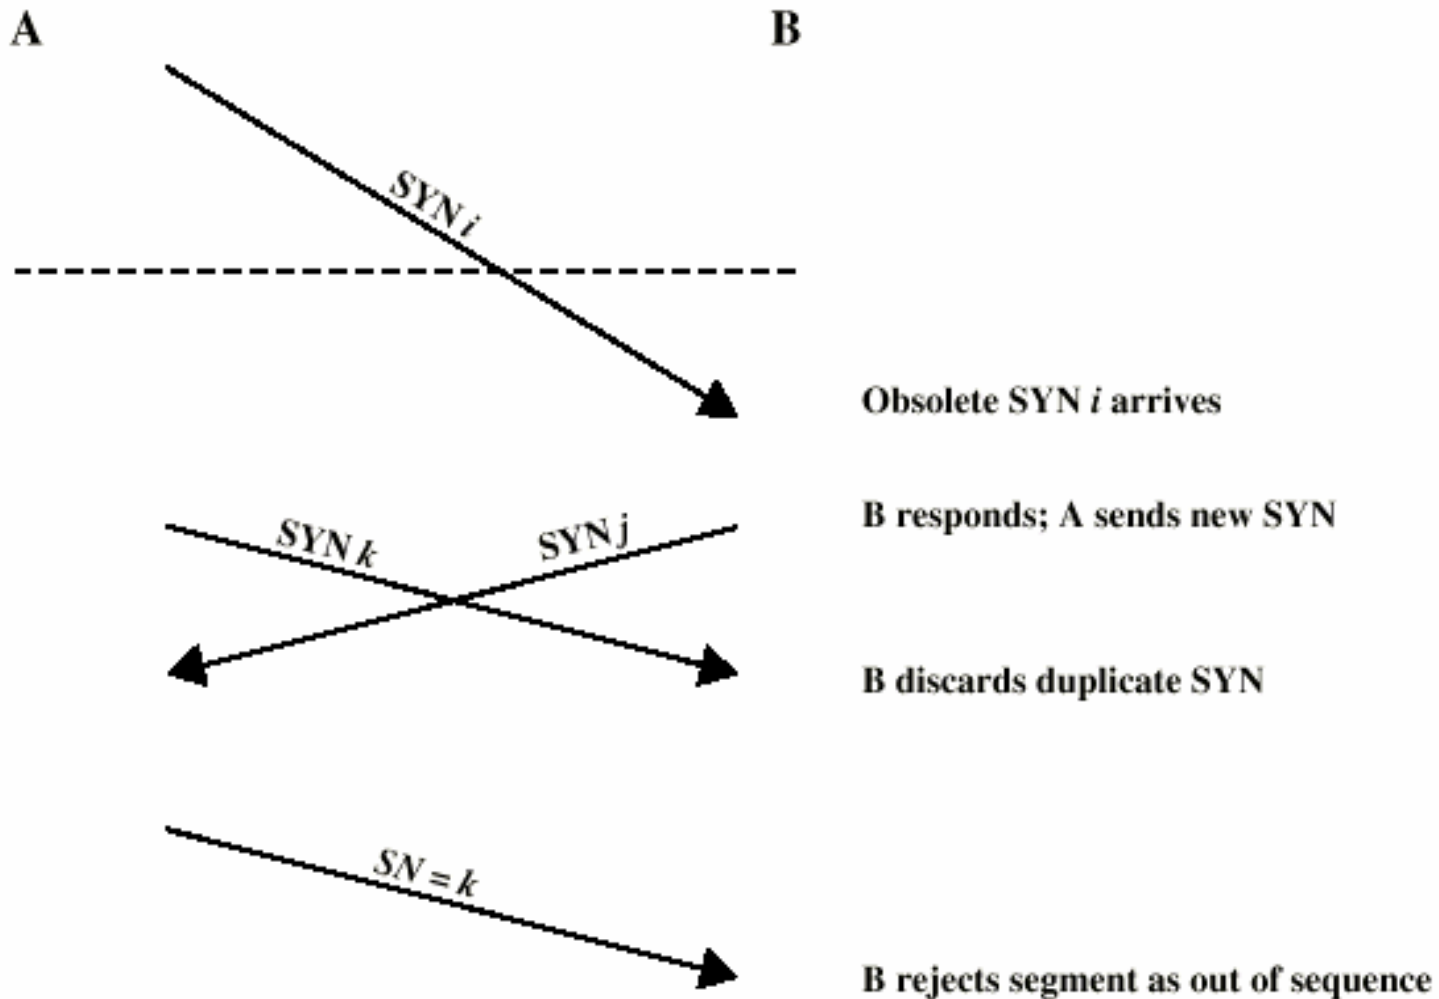
- Two way handshake
 - A send SYN, B replies with SYN
 - Lost SYN handled by re-transmission
 - Can lead to duplicate SYNs
 - Ignore duplicate SYNs once connected
- Lost or delayed data segments can cause connection problems
 - Segment from old connections
 - Start segment numbers are removed from previous connection
 - Use SYN i
 - Need ACK to include i
 - Three Way Handshake



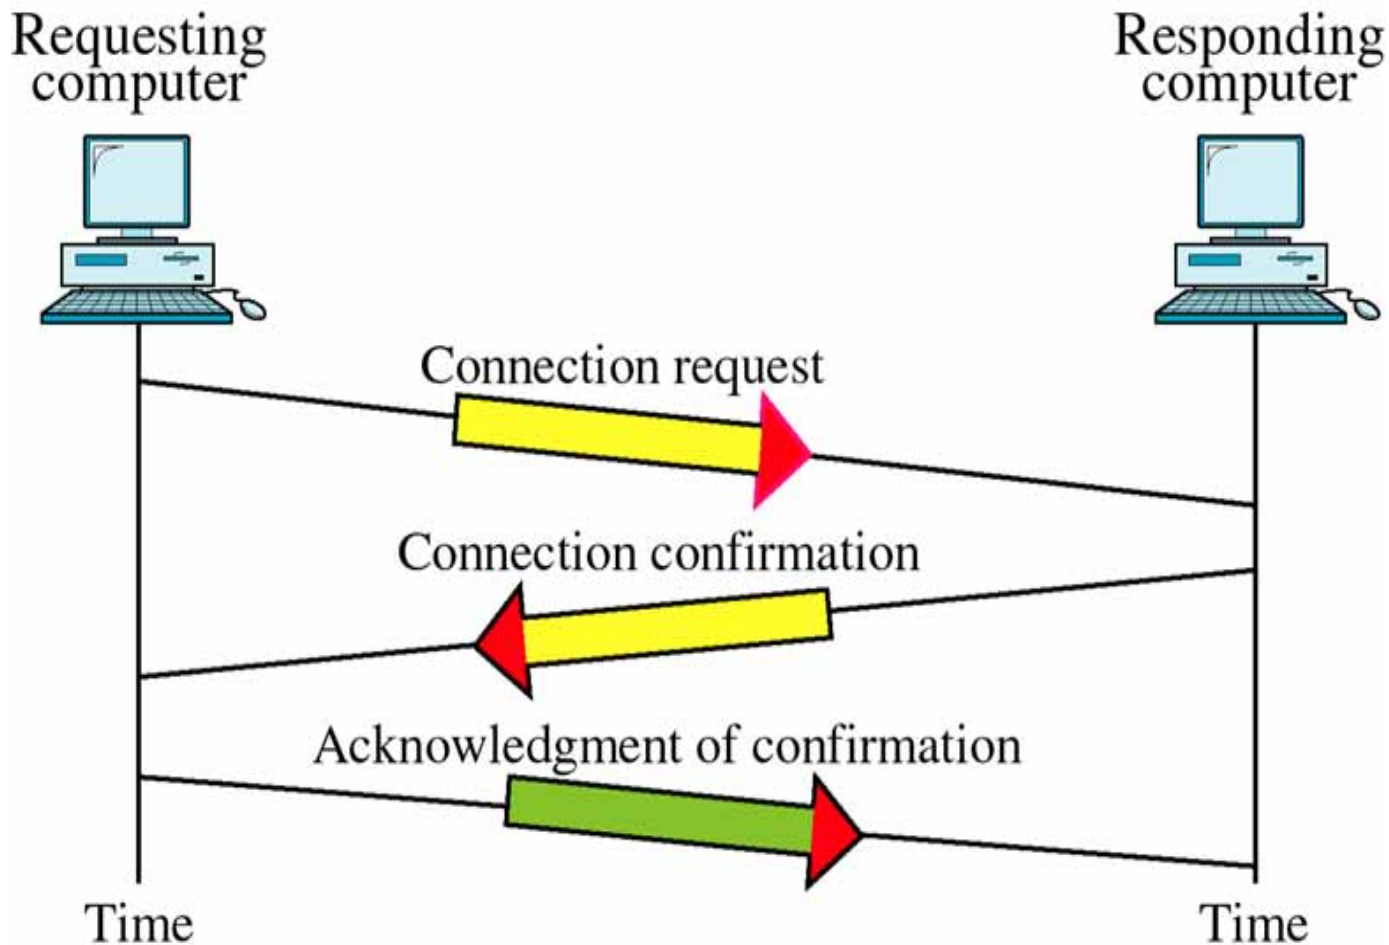


Two Way Handshake: Obsolete Data Segment

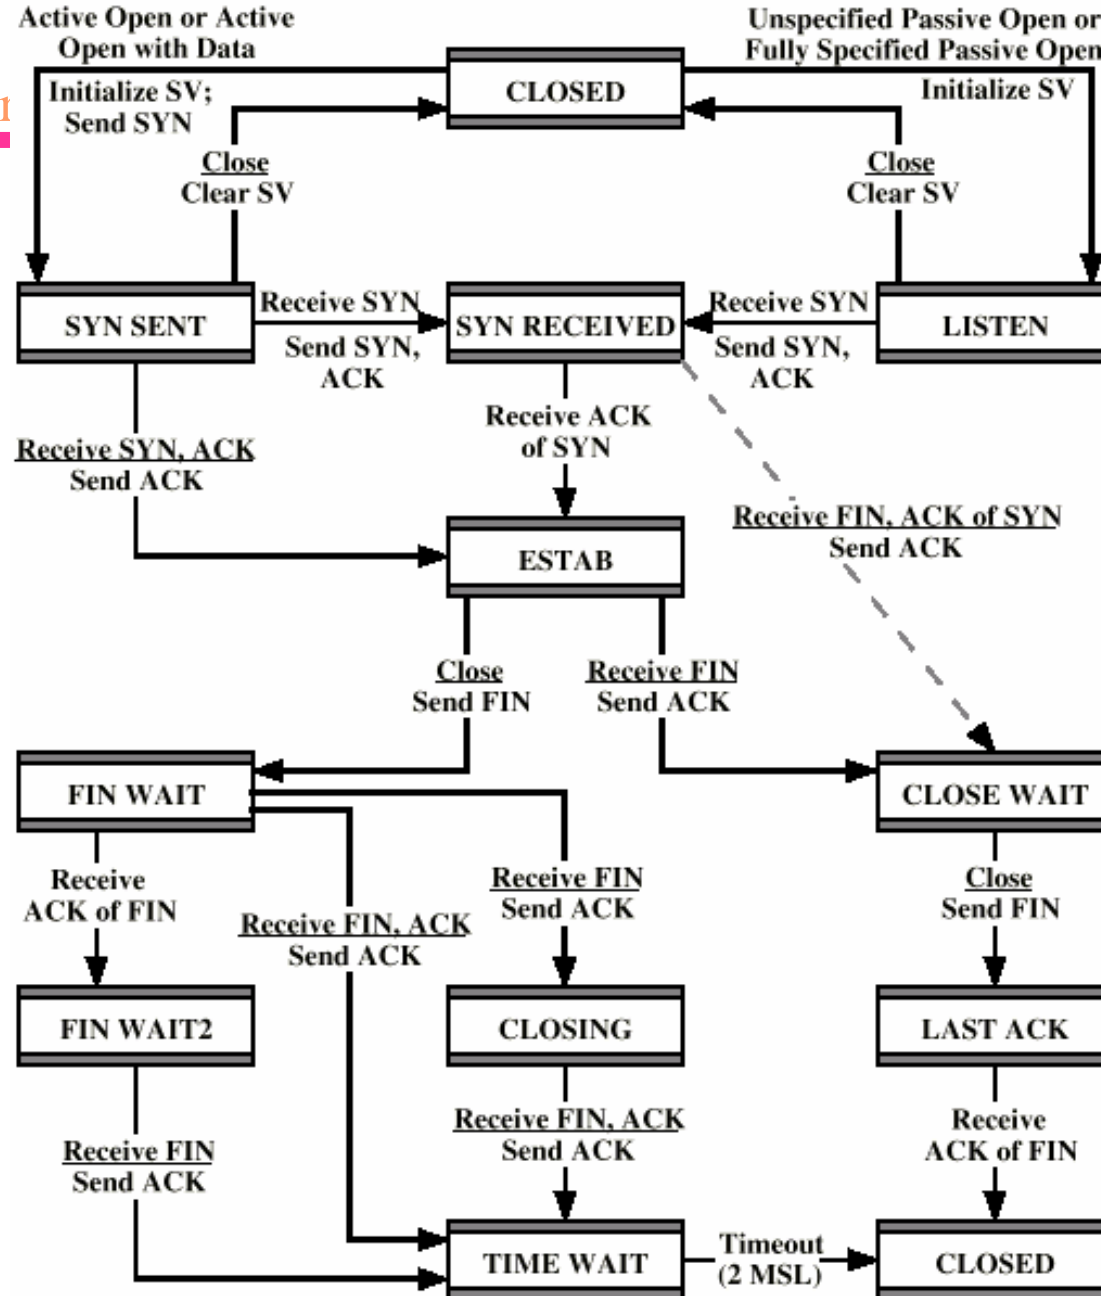
Two Way Handshake: Obsolete SYN Segment



Three Way Handshake: Connection Establishment



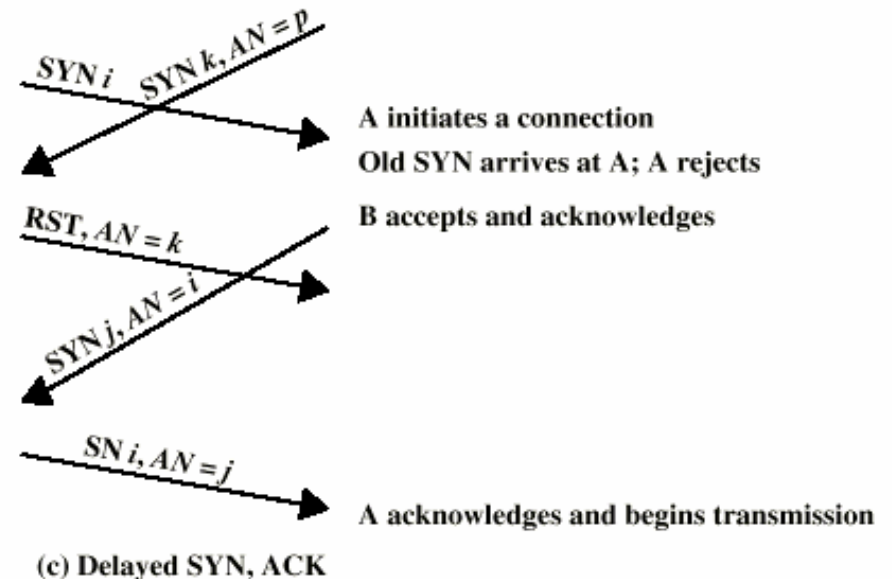

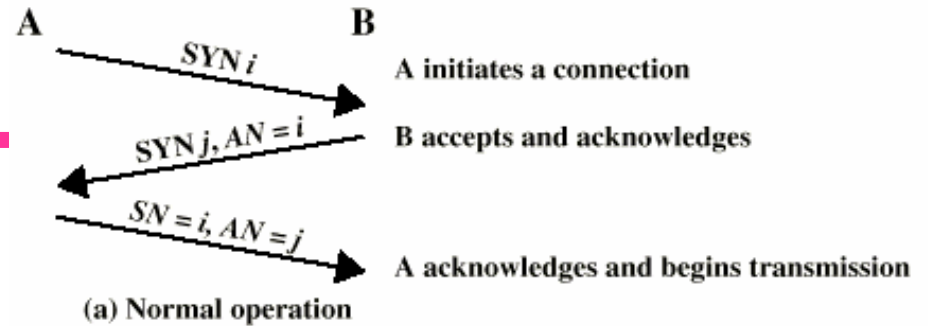
Three Way Handshake: State Diagram



SV = state vector
MSL = maximum segment lifetime



Three Way Handshake: Examples

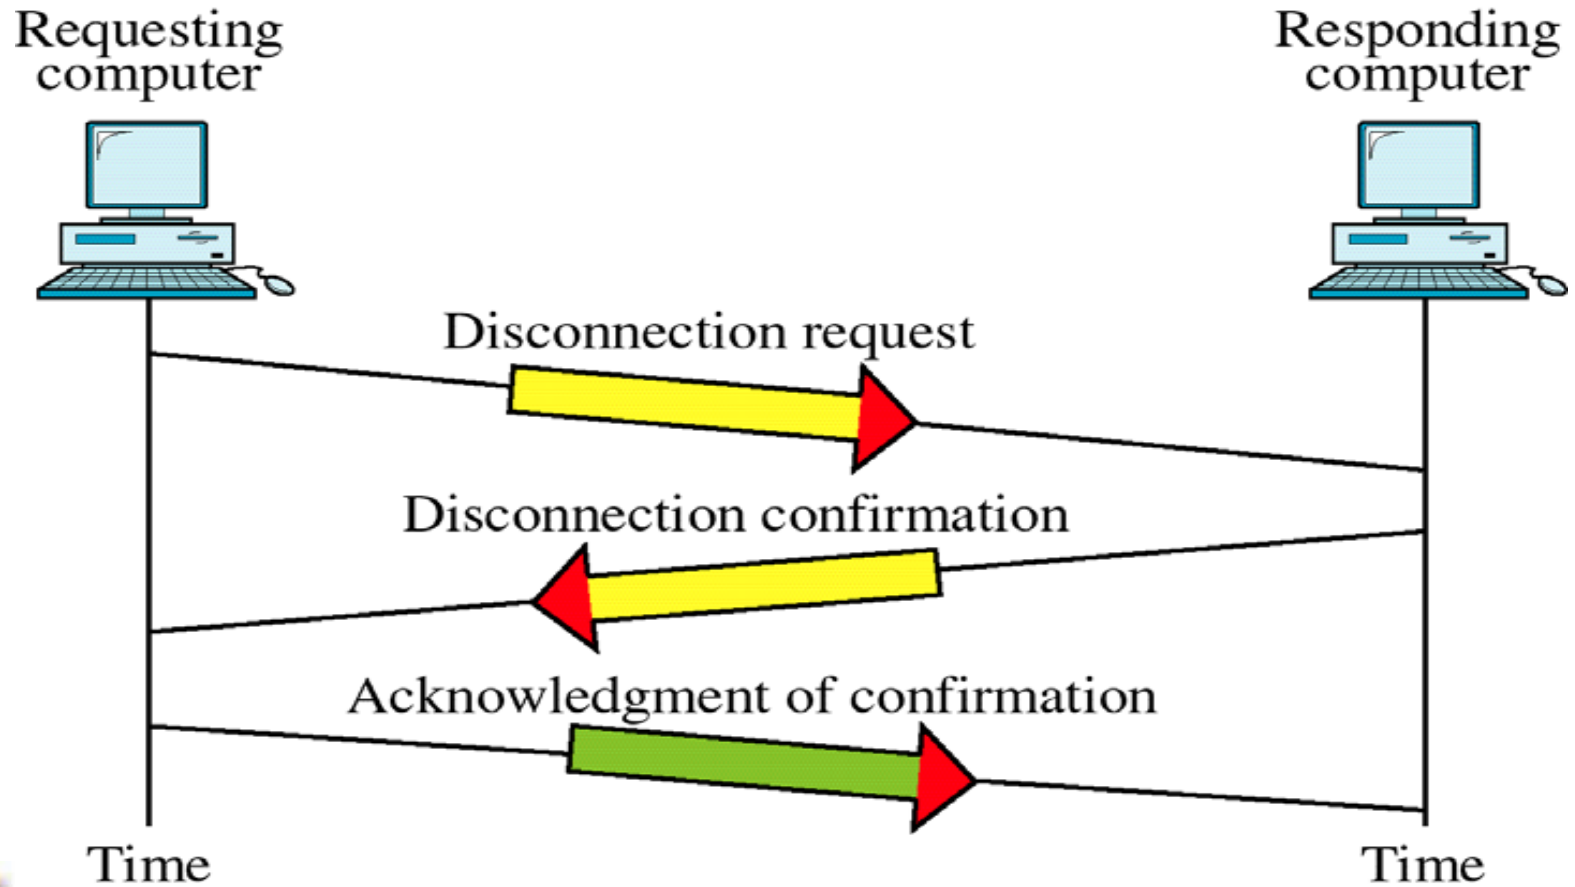


Connection Termination

- Entity in CLOSE WAIT state sends last data segment, followed by FIN
- FIN arrives before last data segment
- Receiver accepts FIN
 - Closes connection
 - Loses last data segment
- Associate sequence number with FIN
- Receiver waits for all segments before FIN sequence number
- Loss of segments and obsolete segments
 - Must explicitly ACK FIN



Connection Termination



Graceful Close

- Send FIN i and receive AN i
- Receive FIN j and send AN j
- Wait twice maximum expected segment lifetime



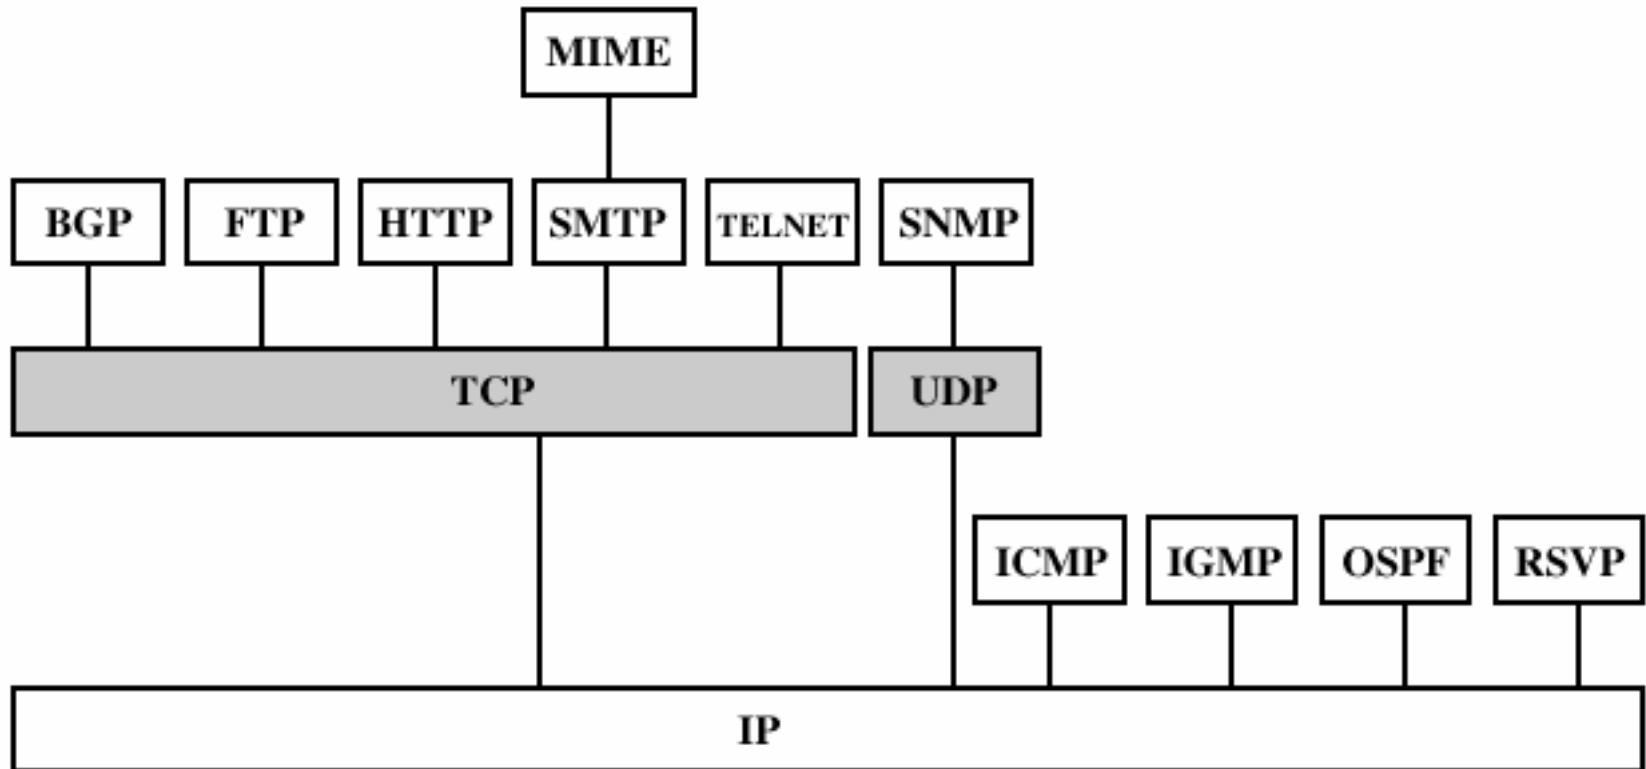
Crash Recovery

- After restart all state info is lost
- Connection is half open
 - Side that did not crash still thinks it is connected
- Close connection using persistence timer
 - Wait for ACK for (time out) * (number of retries)
 - When expired, close connection and inform user
- Send RST i in response to any i segment arriving
- User must decide whether to reconnect
 - Problems with lost or duplicate data

“坚持”而非
“放弃”



Transport-Level Protocols in TCP/IP



TCP & UDP

- Transmission Control Protocol
 - Connection oriented
 - RFC 793
- User Datagram Protocol (UDP)
 - Connectionless
 - RFC 768

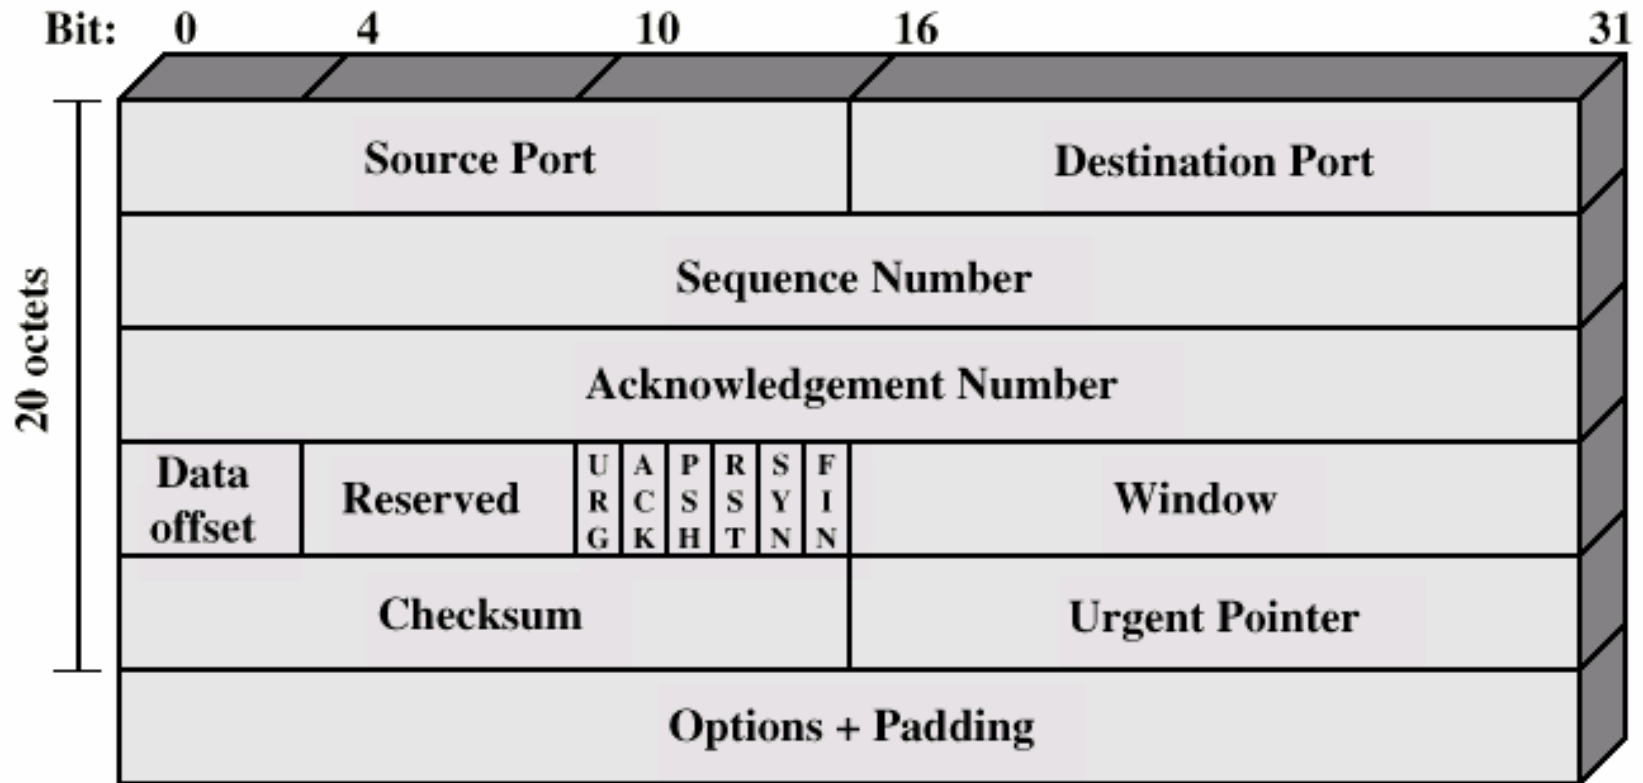


TCP Services

- Reliable communication between pairs of processes
- Across variety of reliable and unreliable networks and internets
- Two labeling facilities
 - Data stream push
 - TCP user can require transmission of all data up to push flag
 - Receiver will deliver in same manner
 - Avoids waiting for full buffers
 - Urgent data signal
 - Indicates urgent data is upcoming in stream
 - User decides how to handle it



TCP Header



Items Passed to IP

- TCP passes some parameters down to IP
 - Precedence
 - Normal delay/low delay
 - Normal throughput/high throughput
 - Normal reliability/high reliability
 - Security



TCP Mechanisms (1)

- Connection establishment
 - Three way handshake
 - Between pairs of ports
 - One port can connect to multiple destinations



TCP Mechanisms (2)

- Data transfer
 - Logical stream of octets
 - Octets numbered modulo 2^{32}
 - Flow control by credit allocation of number of octets
 - Data buffered at transmitter and receiver



TCP Mechanisms (3)

- Connection termination
 - Graceful close
 - TCP users issues CLOSE primitive
 - Transport entity sets FIN flag on last segment sent
 - Abrupt termination by ABORT primitive
 - Entity abandons all attempts to send or receive data
 - RST segment transmitted



Implementation Policy Options

- Send
- Deliver
- Accept
- Retransmit
- Acknowledge



Send

- If no push or close TCP entity transmits at its own convenience
- Data buffered at transmit buffer
 - May construct segment per data batch
 - May wait for certain amount of data



Deliver

- In absence of push, deliver data at own convenience
 - May deliver as each in order segment received
 - May buffer data from more than one segment



Accept

- Segments may arrive out of order
 - In order
 - Only accept segments in order
 - Discard out of order segments
 - In windows
 - Accept all segments within receive window



Retransmit

- TCP maintains queue of segments transmitted but not acknowledged
- TCP will retransmit if not ACKed in given time
 - First only
 - Batch
 - Individual



Acknowledgement

- Immediate
- Cumulative



Congestion Control

- RFC 1122, Requirements for Internet hosts
- Retransmission timer management
 - Estimate round trip delay by observing pattern of delay
 - Set time to value somewhat greater than estimate
 - Simple average
 - Exponential average
 - RTT Variance Estimation (Jacobson's algorithm)



Simple Average

$$ARTT(k+1) = \frac{1}{k+1} \sum_{i=1}^{k+1} RTT(i)$$

- $RTT(i)$ is the round-trip time observed for i th transmission segment
- $ARTT(k)$ is the average round-trip time for the first k segment

$$ARTT(k+1) = \frac{k}{k+1} ARTT(k) + \frac{1}{k+1} RTT(k+1)$$

- With this formulation, it is not necessary to recalculate the entire summation each time



Exponential Average

- Let α replace $\frac{k}{k+1}$, $(0 < \alpha < 1)$
- Smoothed Round Trip Time

$$SRTT(k+1) = \alpha \times SRTT(k) + (1-\alpha) \times RTT(k+1)$$


- Consider the following expansion of above equation:

$$SRTT(k+1) = (1-\alpha)RTT(k+1) + \alpha(1-\alpha)RTT(k) + \\ \alpha^2(1-\alpha)RTT(k-1) + \dots + \alpha^k(1-\alpha)RTT(1)$$




- The older the observation, the less it is counted in the average

Use of Exponential Averaging



(a) Increasing function



(b) Decreasing function



RTO - Retransmission TimeOut

- RTO is transmission timer
 - $RTO(k+1) = SRTT(k+1) + \Delta$
 - $RTO(k+1) = \text{MIN}(\text{MAX}(\beta \times SRTT(k+1), \text{LBOUND}), \text{UBOUND})$
 - $RTO(k+1) = SRTT(k+1) + f \times SDEV(k+1)$



RTT Variance Estimation

- Jacobson's algorithm

$$SRTT(k+1) = (1-g) \times SRTT(k) + g \times RTT(k+1) \quad (g = 1-\alpha)$$

$$SERR(k+1) = RTT(k+1) - SRTT(k)$$

$$SDEV(k+1) = (1-h) \times SDEV(k) + h \times |SERR(k+1)|$$


$$g = 1/8 = 0.125$$

$$h = 1/4 = 0.25$$


$$f = 4$$



Jacobson's RTO Calculation



(a) Increasing function



(b) Decreasing function



Exponential RTO Backoff

- Since timeout is probably due to congestion (dropped packet or long round trip), maintaining RTO is not good idea
- RTO increased each time a segment is re-transmitted
- $RTO = q * RTO$
- Commonly $q=2$
 - Binary exponential backoff



Karn's Algorithm

- If a segment is re-transmitted, the ACK arriving may be:
 - For the first copy of the segment
 - RTT longer than expected
 - For second copy
 - RTT will be much too small
- No way to tell
- Do not measure RTT for re-transmitted segments
- Calculate backoff when re-transmission occurs
- Use backoff RTO until ACK arrives for segment that has not been re-transmitted



Window Management

- Slow start
 - $awnd = \text{MIN}[\text{credit}, cwnd]$
 - Start connection with $cwnd=1$
 - Increment $cwnd$ at each ACK, to some max
- Dynamic windows sizing on congestion
 - When a timeout occurs
 - Set slow start threshold to half current congestion window
 - $ssthresh = cwnd/2$
 - Set $cwnd = 1$ and slow start until $cwnd = ssthresh$
 - Increasing $cwnd$ by 1 for every ACK
 - For $cwnd \geq ssthresh$, increase $cwnd$ by 1 for each RTT






Figure 17.14 Illustration of Slow Start and Congestion Avoidance



UDP

- User datagram protocol
- RFC 768
- Connectionless service for application level procedures
 - Unreliable
 - Delivery and duplication control not guaranteed
- Reduced overhead
- e.g. network management (Chapter 19)




UDP Uses

- Inward data collection
- Outward data dissemination
- Request-Response
- Real time application



UDP Header



Required Reading

- Stallings, D&CC, 6e, Chapter 17
- Forouzan, Introduction to DC&N, 2e, Chapter 22
- 谢希仁, 计算机网络, 第三版, 第八章
- RFC 793 ,768, 1122, 2001



Problems

- 17.3, 17.9, 17.15, 17.20

